


# CS161: Operating Systems

Matt Welsh  
mdw@eecs.harvard.edu



Lecture 2: OS Structure and System Calls  
February 6, 2007

# Lecture Overview

## Protection Boundaries and Privilege Levels

- What makes the kernel different from regular user programs

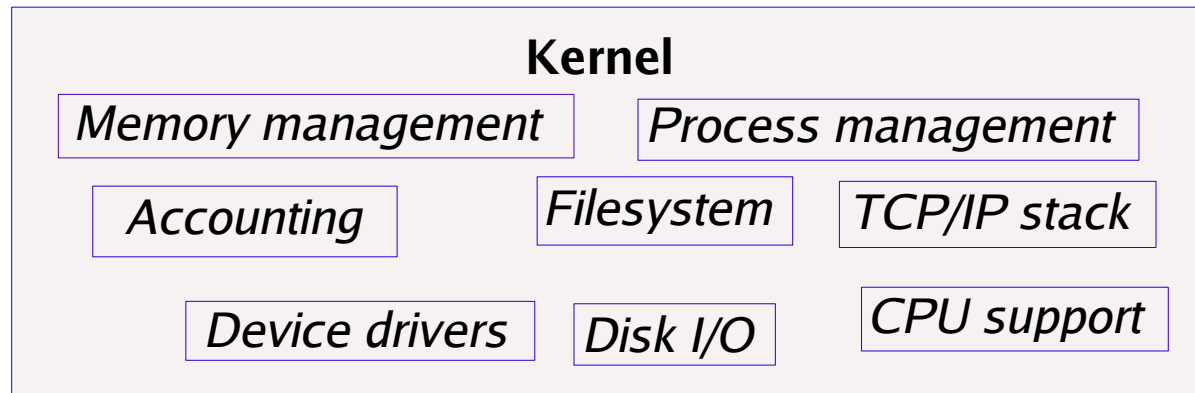

## System calls and hardware interrupts

- How user applications and the kernel interact with each other

## Operating System Structure

- What are the major functional components of an OS?
- How are those components interrelated?
- What are the interfaces to those components?

# Operating System Overview



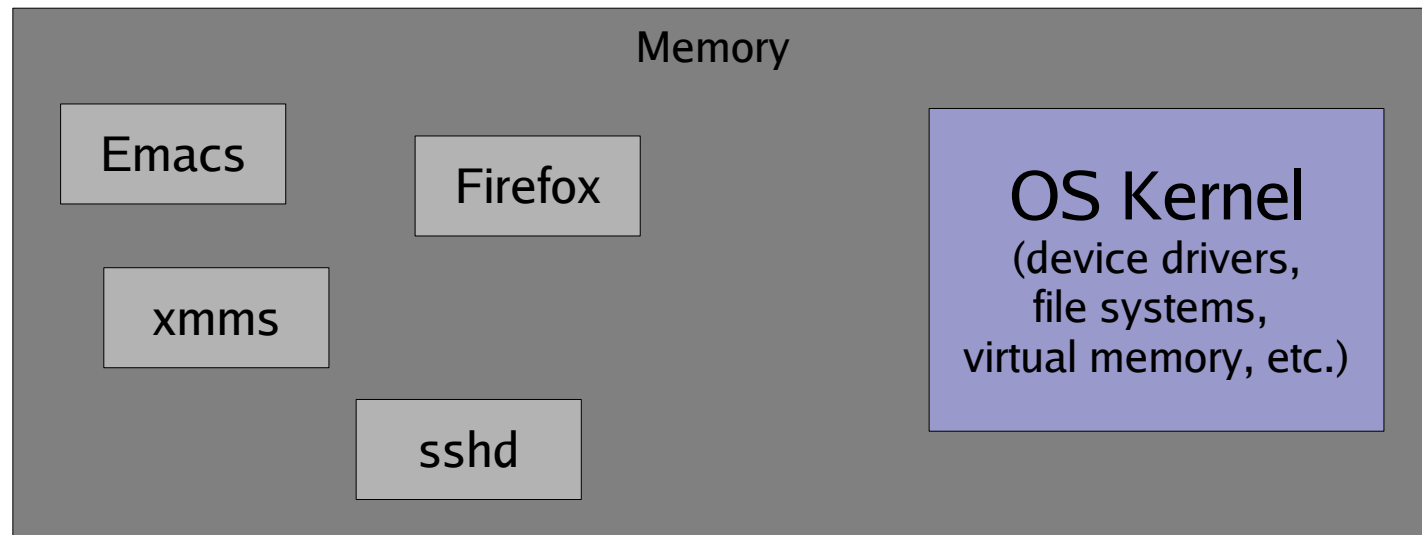
Hardware/software interface



*Gnarly world of hardware*

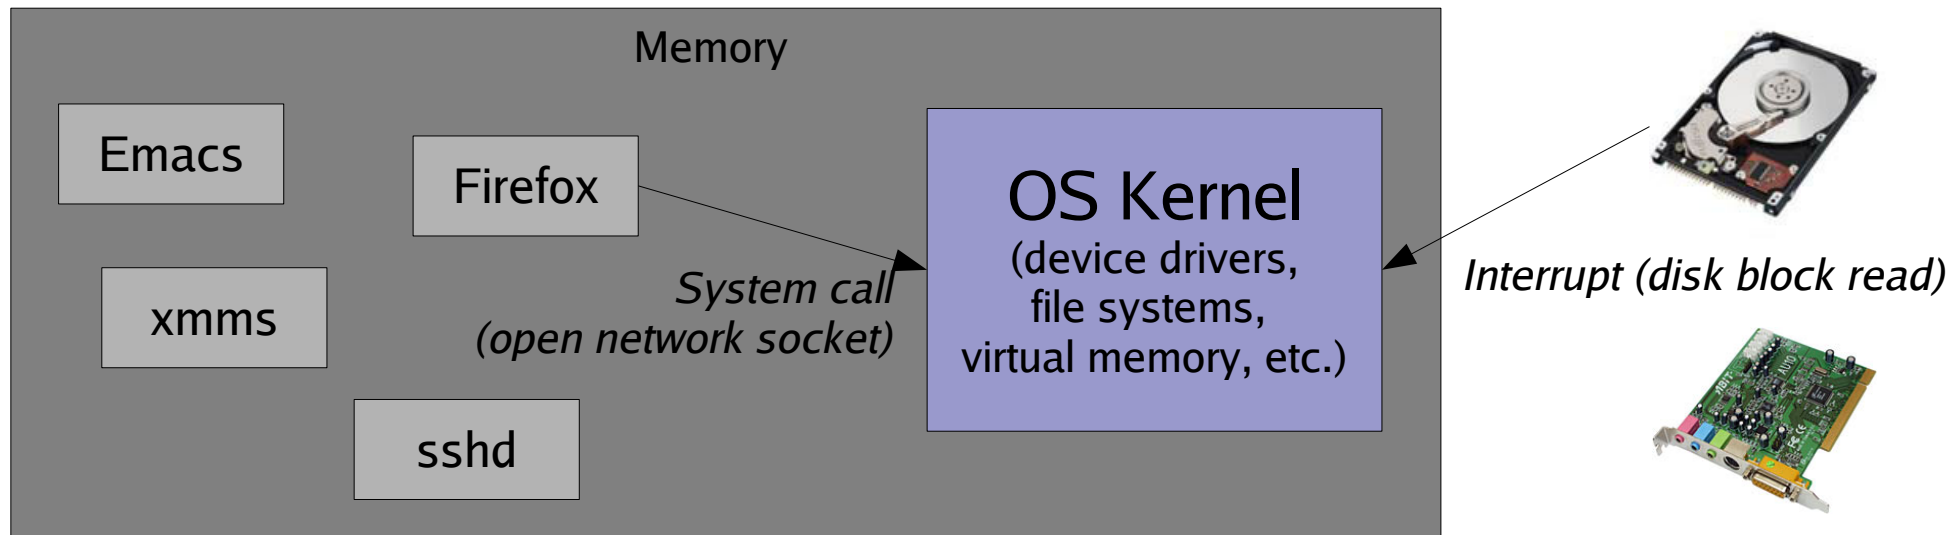
# Operating System basics

The OS kernel is just a bunch of code that sits around in memory, waiting to be executed



# Operating System basics

The OS kernel is just a bunch of code that sits around in memory, waiting to be executed



OS is triggered in two ways: *system calls* and *hardware interrupts*

System call: Direct “call” from a user program

- For example, `open()` to open a file, or `exec()` to run a new program

Hardware interrupt: Trigger from some hardware device

- For example, when a disk block has been read or written

*How else might the kernel get control ???*

# Interrupts – a primer

An *interrupt* is a signal that causes the CPU to jump to a pre-defined instruction – called the *interrupt handler*

- Interrupt can be caused by hardware or software

## Hardware interrupt examples

- Timer interrupt (periodic “tick” from a programmable timer)
- Device interrupts
  - *e.g., Disk will interrupt the CPU when an I/O operation has completed*

## Software interrupt examples (also called *exceptions*)

- Division by zero error
- Access to a bad memory address
- Intentional software interrupt – e.g., x86 “INT” instruction
  - *Can be used to trap from user program into the OS kernel!*
  - *Why might this be useful?*

# Interrupt handler example

1) OS fills in interrupt handler table (usually at boot time)


Interrupt handler table



Interrupt handler  
for interrupt 4

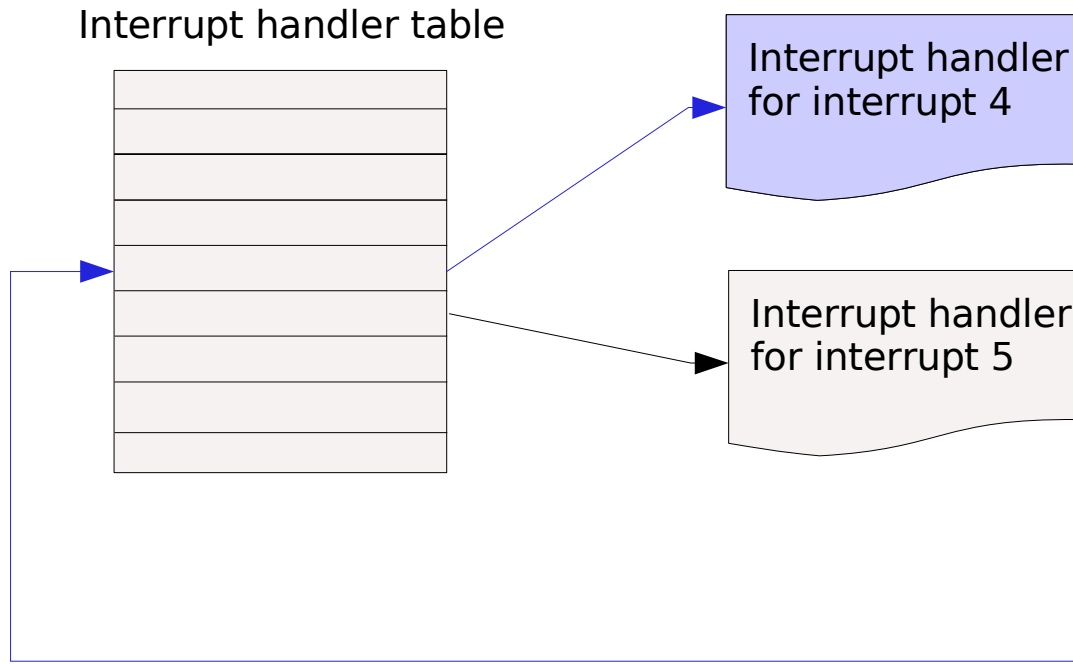
Interrupt handler  
for interrupt 5

2) Interrupt occurs – e.g., hardware signal



3) CPU state saved to stack

# Interrupt handler example



1) OS fills in interrupt handler table (usually at boot time)

2) Interrupt occurs – e.g., hardware signal



3) CPU state saved to stack

4) CPU consults interrupt table and invokes appropriate handler



# Protection

A major job of the OS is to enforce *protection*

Prevent malicious or buggy programs from:

- Allocating too many resources (denial of service)
- Corrupting or overwriting shared resources (files, shared memory, etc.)

Prevent different users, groups, etc. from:

- Accessing or modifying private state (files, shared memory, etc.)
- Killing each other's processes

A lot of viruses, worms, etc. exploit security holes in the OS

- Overrunning a memory buffer in the kernel can give a non-root process root privileges
  - *Kernel code needs to be rock solid in order to be secure!!!*

How does the OS enforce protection boundaries?

# User mode vs. kernel mode

What makes the kernel different from user programs?

- Kernel can execute special *privileged instructions*

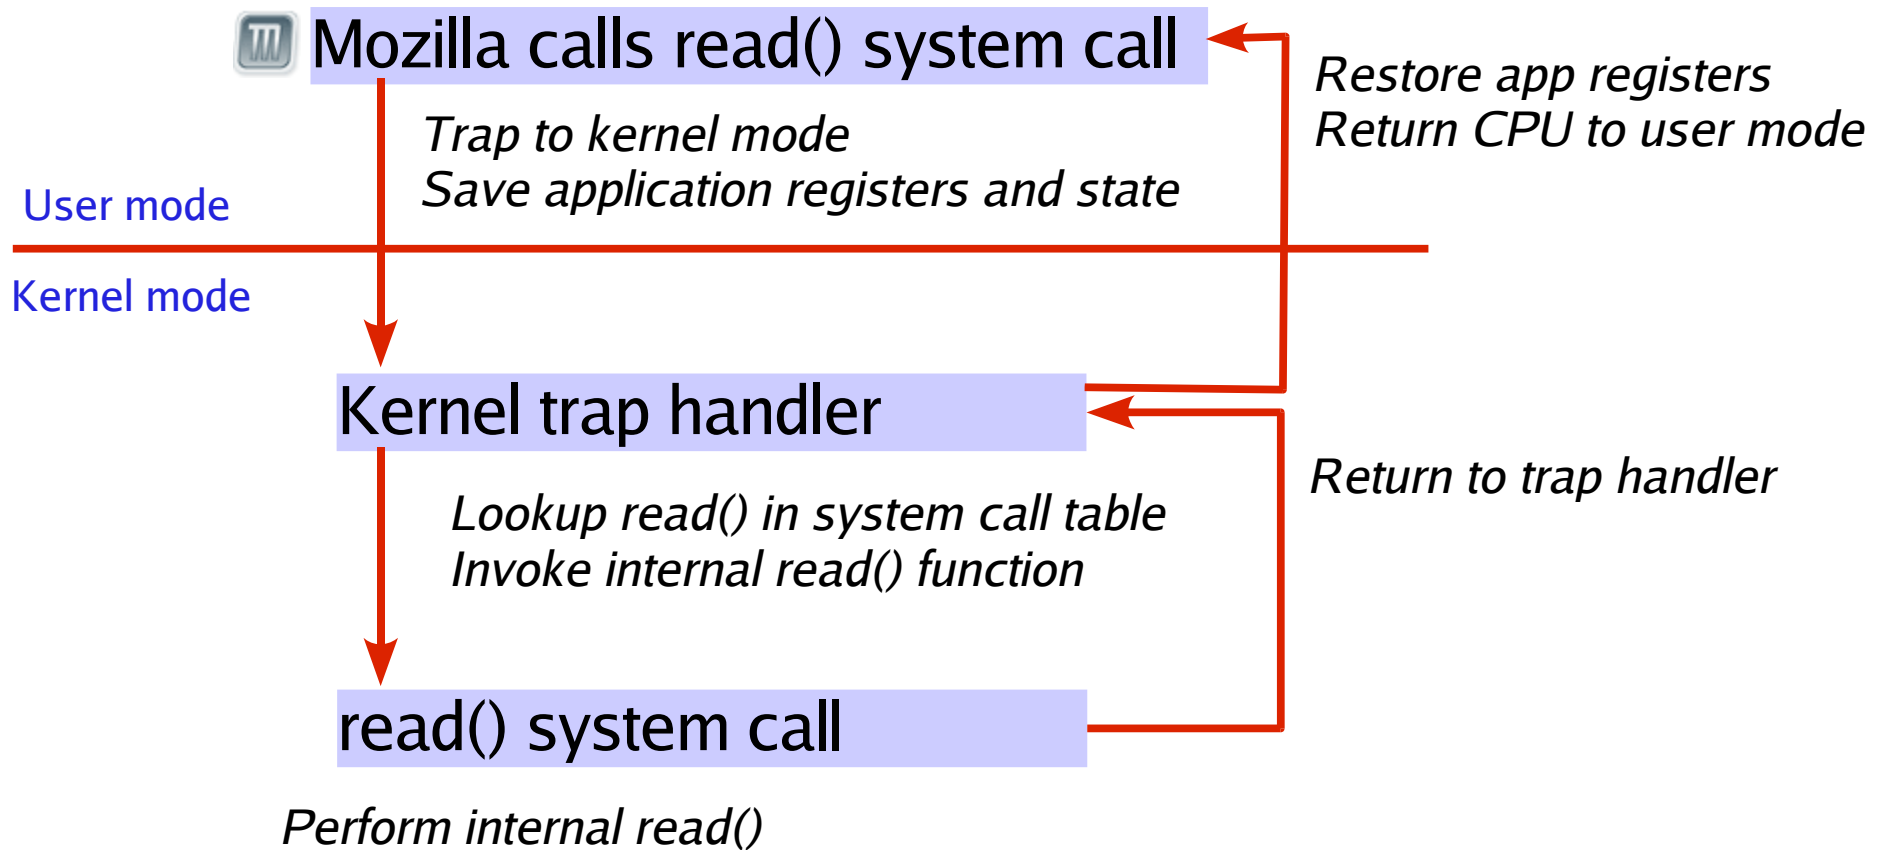
Examples of privileged instructions:

- Access I/O devices
  - *Poll for IO, perform DMA, catch hardware interrupt*
- Manipulate memory management
  - *Set up page tables, load/flush the TLB and CPU caches, etc.*
- Configure various “mode bits”
  - *Interrupt priority level, software trap vectors, etc.*
- Call halt instruction
  - *Put CPU into low-power or idle state until next interrupt*

These are enforced by the **CPU hardware itself**.

- CPU has at least two protection levels: *Kernel mode* and *user mode*
- CPU checks current protection level on each instruction
- What happens if user program tries to execute a privileged instruction?

# Boundary Crossing



# Protection Rings

On most CPU designs, there are just two protection levels:

- Kernel mode and user mode

Some CPUs, however, support multiple protection levels


The Intel x86 family has four protection “rings”

Ring 0 (innermost ring) has the OS kernel

Ring 3 has application code

Rings 1 and 2 may be used for less-privileged OS code

- *e.g., Third-party drivers*



# Intel x86 Protection Ring Rules

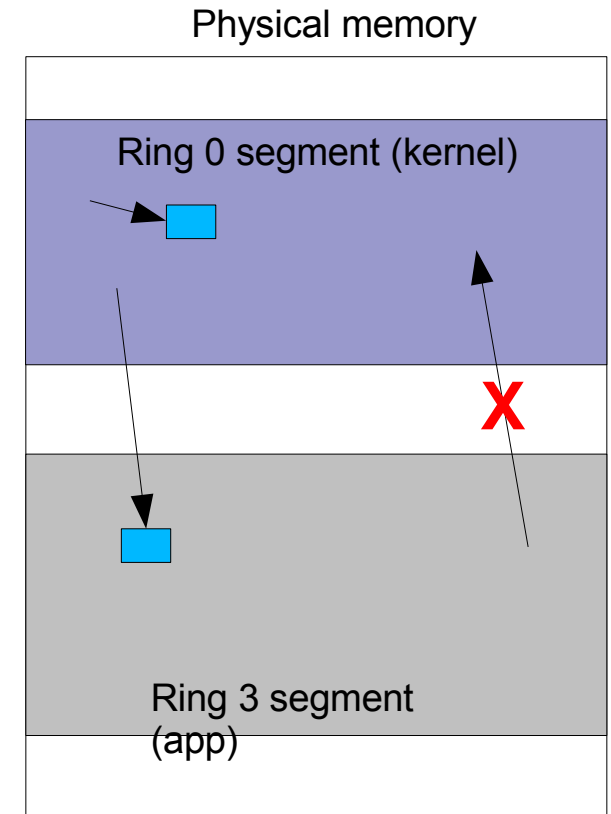
Each memory segment has an associated privilege level, 0 through 3

The CPU has a Current Protection Level (CPL)

- Usually the privilege level of the segment where the program's instructions are being read from

Program can read/write data in segments of *lower privilege* than CPL.

- e.g., Kernel can read/write user memory
- But, user cannot read/write kernel memory
  - (Why?)



# Intel x86 Protection Ring Rules

Each memory segment has an associated privilege level, 0 through 3

The CPU has a Current Protection Level (CPL)

- Usually the privilege level of the segment where the program's instructions are being read from

Program can read/write data in segments of *lower privilege* than CPL.

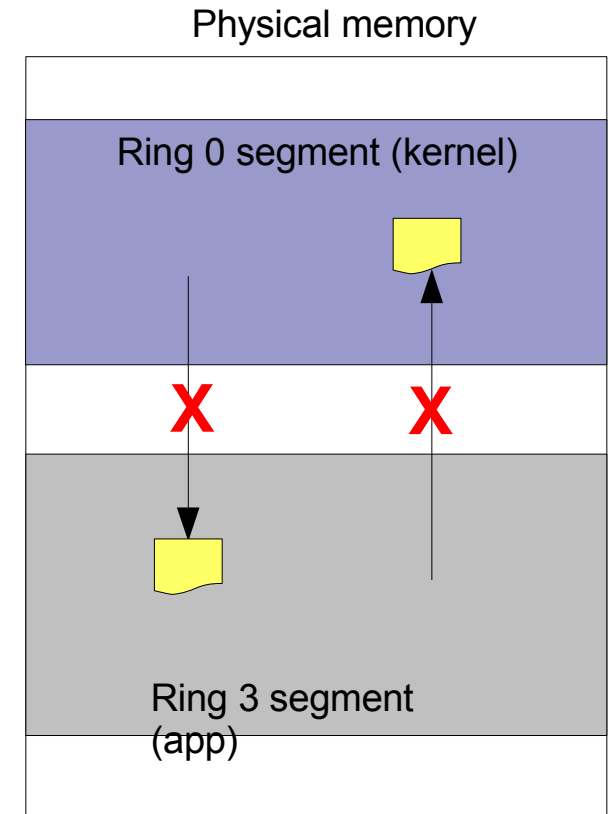
- e.g., Kernel can read/write user memory
- But, user cannot read/write kernel memory

Program cannot (directly) call code in *higher-privilege* segments

- *Why?*

Program cannot (directly) call code in *lower-privilege* segments

- *Why?*



# Intel x86 Protection Ring Rules

A gate is a special memory address that “opens a door” from a low-privilege segment to a high-privilege one.

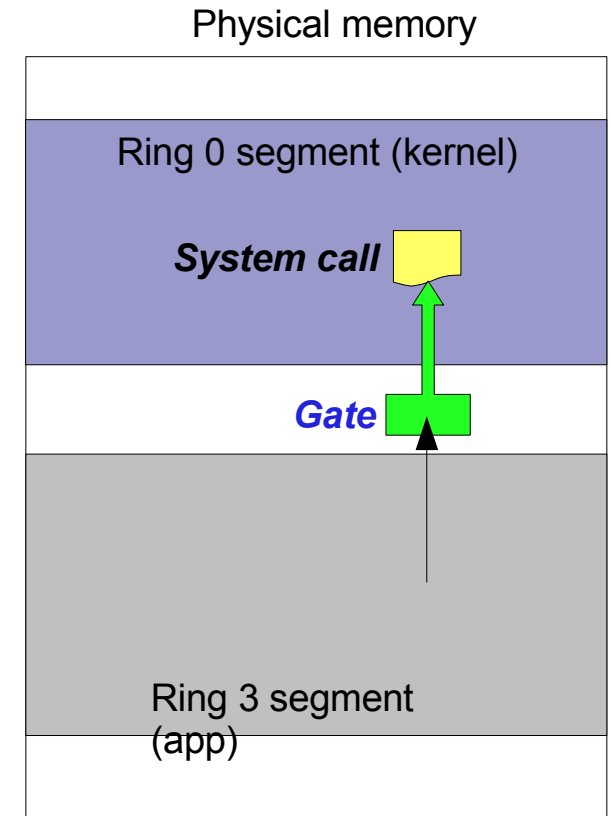
- When a low-privilege program calls a gate, it automatically raises its CPL to the higher level.
- When returning from a gate subroutine, CPL automatically dropped to original level.

Gates allow high-privilege program to control the entry points into its code.

- For example, OS system calls can be implemented using gates.

Gates do not allow higher-privileged code to call lower-privileged code.

- *Why?*



# System Call Issues

Kernel must verify the arguments that it is passed – why?

How does the application pass in data to the kernel?

- Example: `write()` passes in a pointer to a buffer of data to write to the file

How does the OS return kernel state from a system call?

- Example: `read()` returns an int indicating the number of bytes read



# Uniprogramming and Multiprogramming

## Uniprogramming

- Only one program can run at a given time on the system
- Like old batch systems, MS-DOS, etc.

## Multiprogramming (a.k.a. “multitasking”)

- Multiple programs can run simultaneously
- Although only one program running **at any given instant**
  - *(Unless you have multiple CPUs!!!!)*

## Tradeoffs

- Writing a uniprogramming OS is simpler
  - *Why?*
- But, multitasking OSs use resources more efficiently
  - *Why?*

### **Note on terminology:**

**Multitasking/multiprogramming** refer to the number of programs running

**Multiprocessing** refers to the number of CPUs in the system

# Process Management

An OS executes many kinds of applications

- Regular user programs
  - *Emacs, Mozilla, this OpenOffice program, etc...*
- Administrative servers
  - *Cron: Runs jobs at pre-scheduled times*
  - *Sshd: Manages incoming ssh connections*
  - *Lpd: Queues up jobs for the printer*

Each of these activities is encapsulated in a *process*

- A process consists of three main parts:

## Processor state

- registers, program counter

## OS resources

- open files, network sockets, etc.

## Address space:

- The memory that a process accesses – its code, variables, stack, etc.

# Process Example

A **process** is an instance of a program being executed

- Use “ps” to list processes on UNIX systems

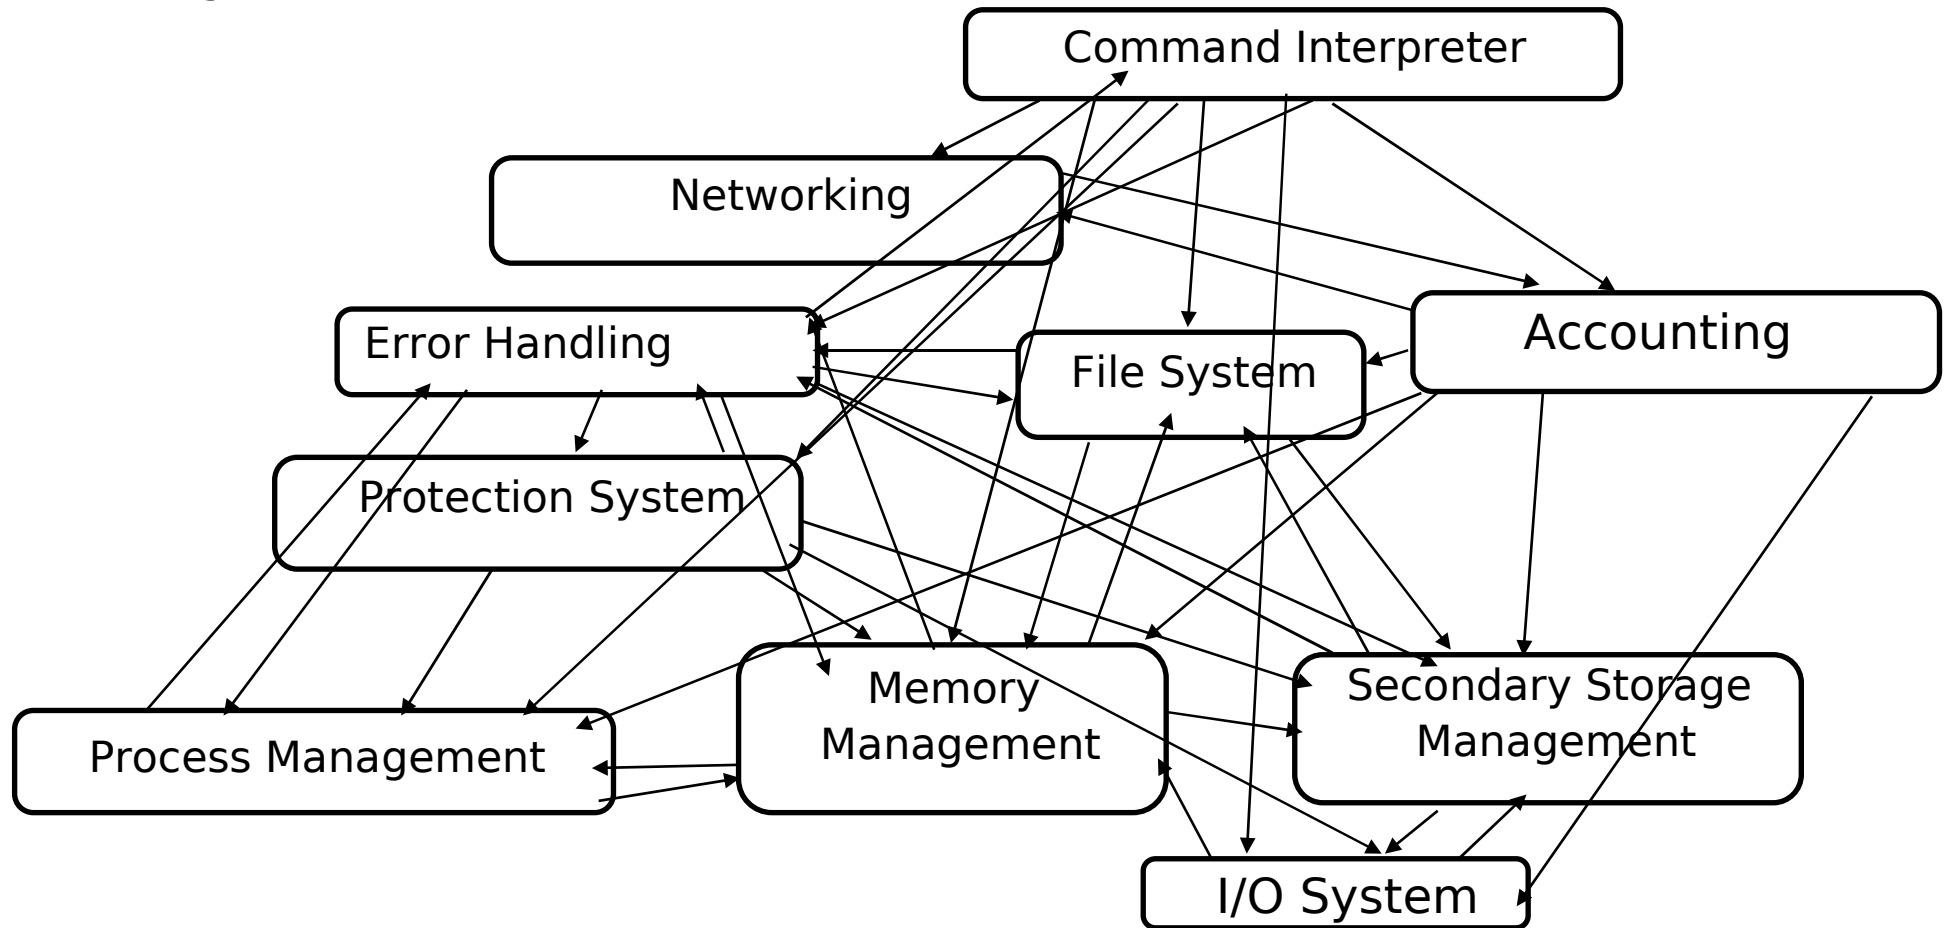
PID	TTY	STAT	TIME	COMMAND
842	tty1	S	0:00	-bash
867	tty1	S	0:00	xinit
873	tty1	S	0:00	fvwm2
887	tty1	S	0:00	xload
888	tty1	S	0:02	/usr/local/j2sdk1.4.0/bin/java ApmView 896 243
1881	tty1	S	0:00	rxvt -fn fixed -cr red -fg white -bg #586570 -geometr
1883	pts/2	S	0:00	bash
1910	pts/0	S	0:00	/bin/sh /home/mdw/bin/ooffice arch.sxi
1911	pts/0	S	1:20	/usr/local/OpenOffice.org1.1.0/program/soffice.bin ar
1937	tty1	S	0:00	/bin/sh /home/mdw/bin/set-wlan-OFF
2310	pts/2	R	0:00	ps -Umdw -x

↑ terminal  
↑ Status  
↑ Command line  
↑  
Process ID      Total CPU time

The next few lectures will deal with details of process creation and management...

# OS Structure

It's not always clear how to tie these different OS components together...

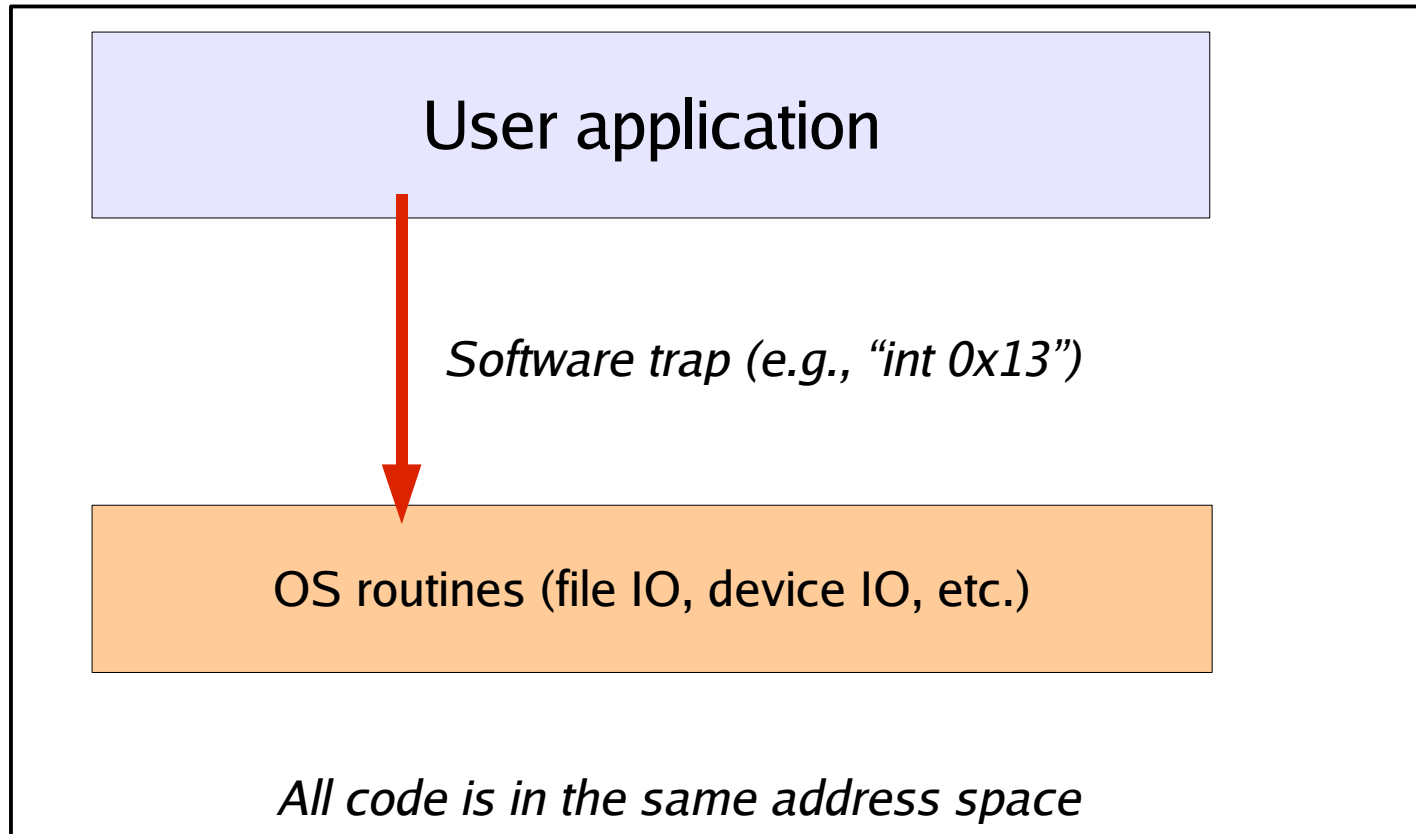


(For illustrative purposes only)

# “Executive” model

Not really an OS “kernel” per se ...

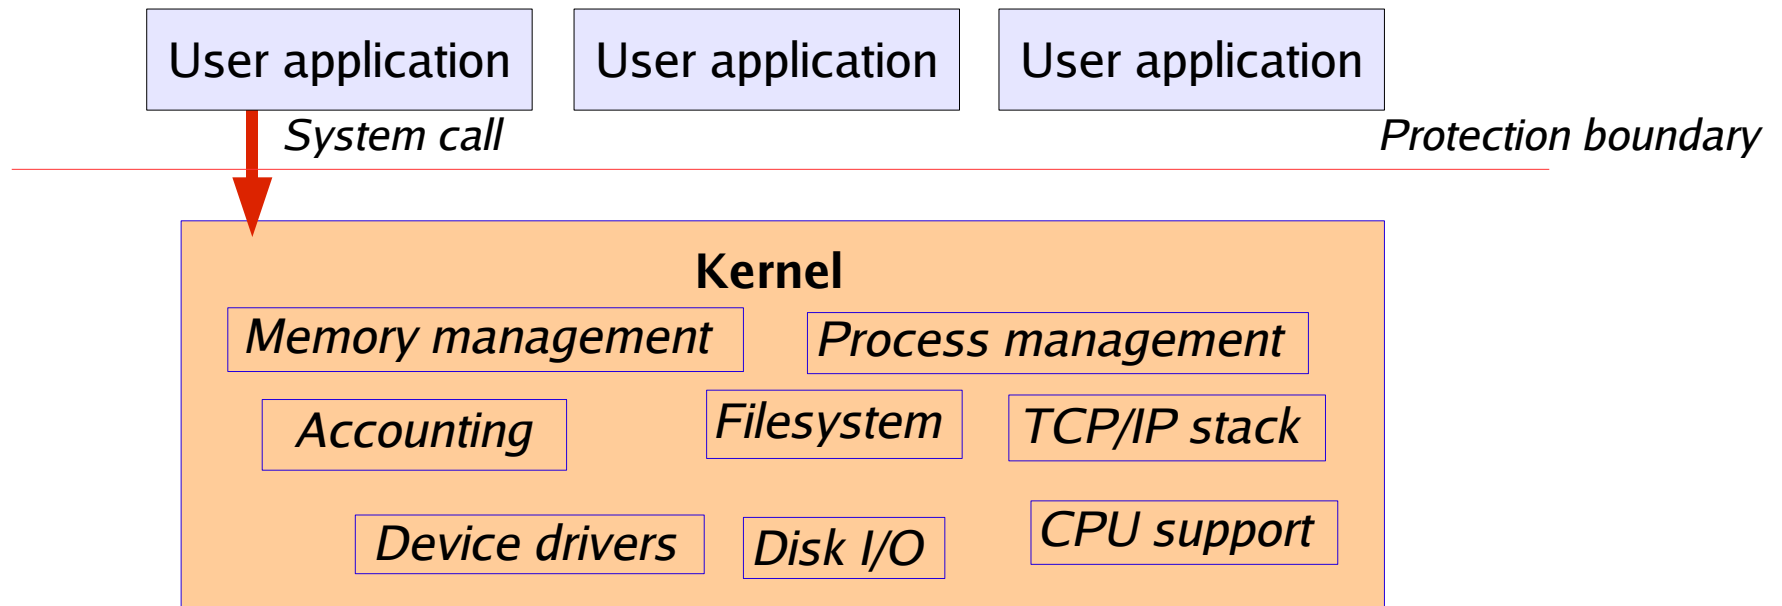
- But rather a collection of routines, resident in memory, that the application can invoke
  - *MS-DOS used this structure*
- Effectively a library, except the application is not explicitly linked to the OS routines
  - *Instead, uses software traps to invoke OS routines from user code*



# Monolithic Kernels

## Most common OS kernel design

- Kernel code is privileged and lives in its own address space
- User applications are unprivileged and live in their own separate address spaces
- All kernel functions loaded into memory as one large, messy program

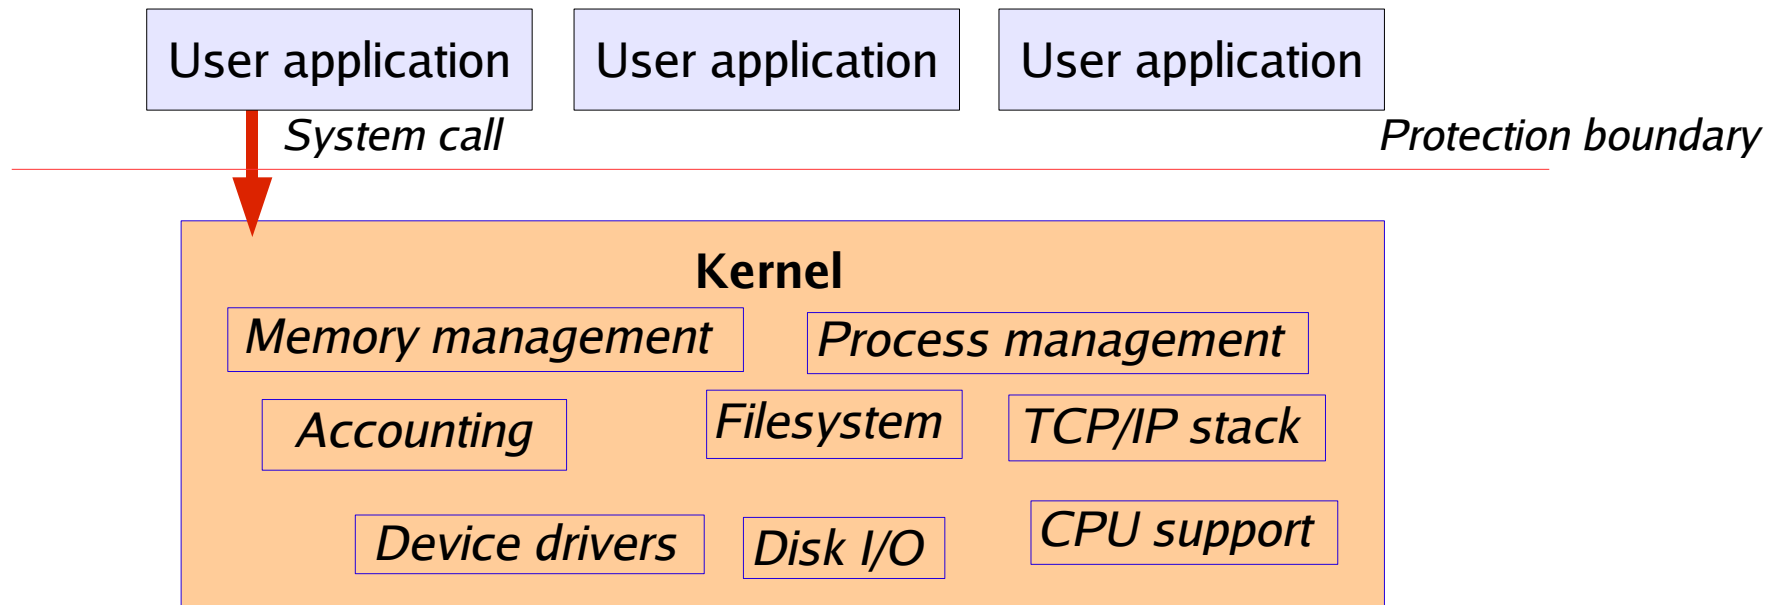


Pros and cons???

# Monolithic Kernels

## Most common OS kernel design

- Kernel code is privileged and lives in its own address space
- User applications are unprivileged and live in their own separate address spaces
- All kernel functions loaded into memory as one large, messy program



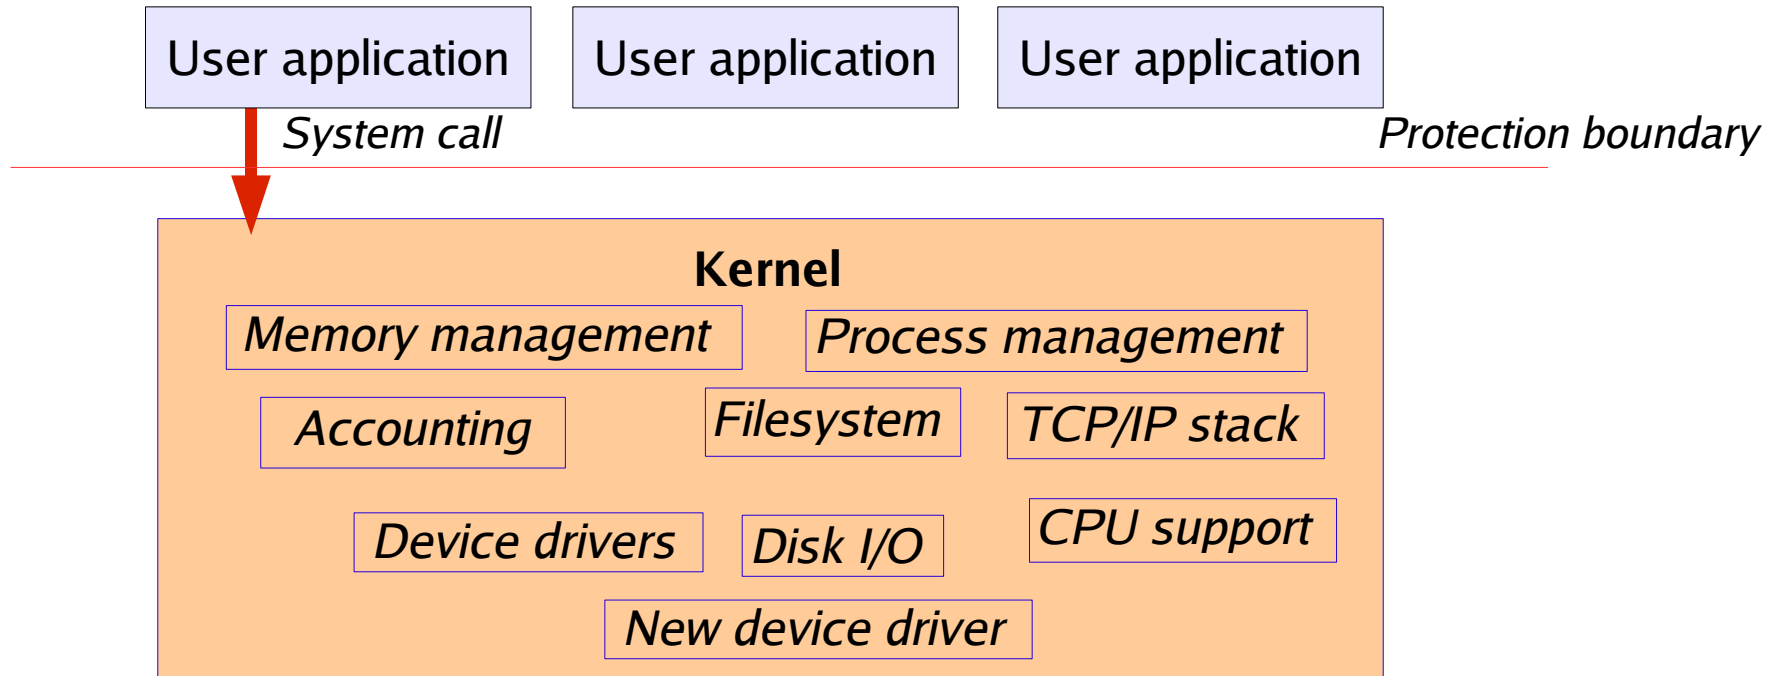
## Pros and cons

- Pro: Overhead of module interactions *within* the kernel is low (function call)
- Pro: Kernel modules can directly share memory
- Con: Very complicated and difficult to organize
- Con: A bug in *any part* of the kernel can crash the whole system!

# Loadable Modules

Allows new kernel code to be dynamically loaded and unloaded

- The kernel is otherwise monolithic, as before




Pros and cons???



# Loadable Modules

Allows new kernel code to be dynamically loaded and unloaded

- The kernel is otherwise monolithic, as before



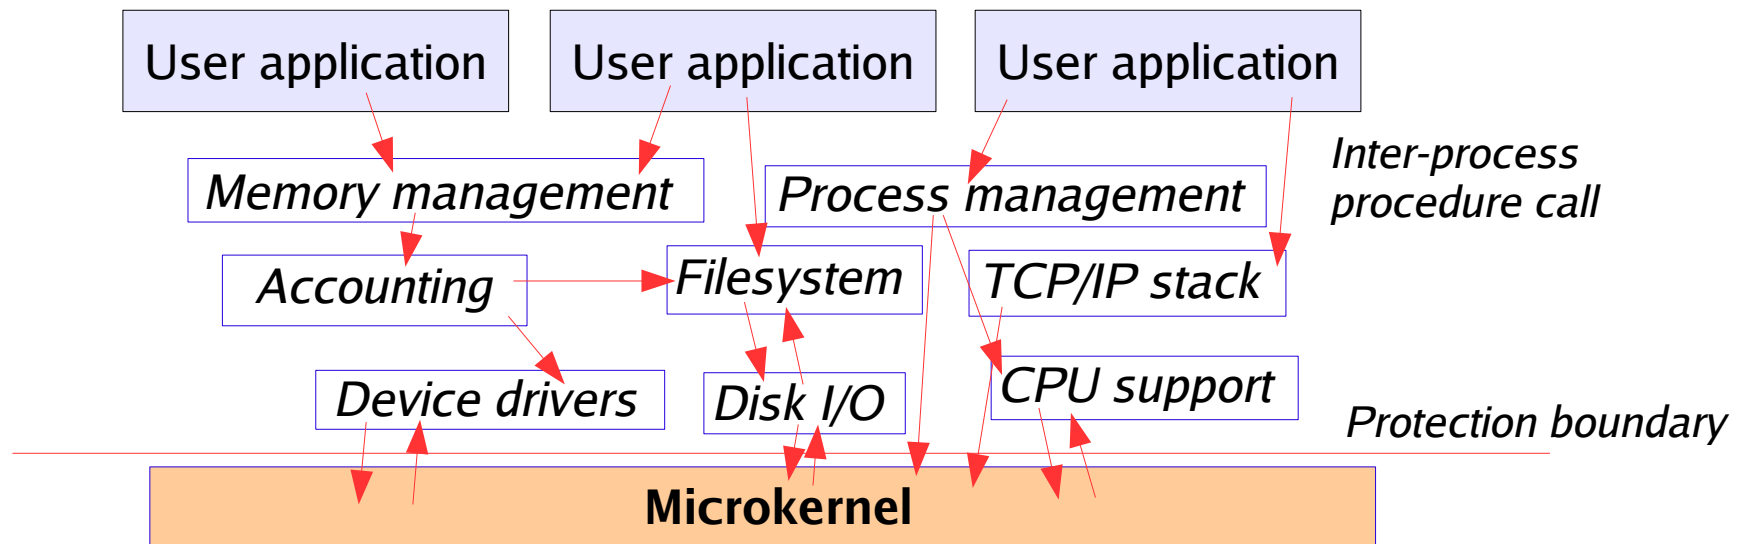
## Pros and cons

- Pro: Save memory by only loading drivers, etc. that are needed
- Pro: Makes it easier to develop new kernel code outside of the main tree
- Con: Once loaded, module can wreak havoc on the running kernel

# Microkernels

Use a very small, minimal kernel, and implement all other functionality as user level “servers”

- Kernel only knows how to perform lowest-level hardware operations
- Device drivers, filesystems, virtual memory, etc. all implemented on top
- Use inter-process procedure call (IPC) to communicate between applications and servers



Pros and Cons???

# Microkernels 2

## Pros and cons

- Pro: Kernel is small and simple, servers are protected from each other
- Con: Overhead of invoking servers may be very high
  - *e.g., A user process accessing a file may require inter-process communication through 2 or 3 servers!*

## Microkernels today

- Lots of research in late 80's and early 90's
- Windows NT uses “modified microkernel”:
  - *Monolithic kernel for most things, OS APIs (DOS, Win3.1, Win32, POSIX) implemented as user-level services*
- Mac OS X has reincarnated the microkernel architecture as well:
  - *Gnarly hybrid of Mach (microkernel) and FreeBSD (monolithic)*



# Administrative Stuff

Next Lecture: Dive into the process subsystem

- How the OS manages loading and running executables

Read Tanenbaum 2.1